

Analytical Design-Space Exploration and Silicon Validation of a Depthwise-Separable CNN Keyword-Spotting Accelerator

Shlok Sharma
Northeastern University
Boston, MA, USA
sharma.shlok@northeastern.edu

Abstract—Always-on keyword spotting (KWS) runs on billions of battery-powered edge devices, which makes per-inference energy a first-order design constraint. Architects routinely prune an accelerator design space with fast analytical models such as Timeloop/Accelergy long before committing to register-transfer level (RTL) code, yet the fidelity of those early estimates against real silicon is seldom quantified end to end on an open, reproducible flow. This paper is a case study that closes that loop for a depthwise-separable CNN (DS-CNN) KWS accelerator. We run an analytical design-space exploration, *commit* quantitative pre-RTL predictions, and then implement, functionally verify, and synthesize the compute core to the open-source SkyWater 130 nm process design kit, comparing the two. The area estimate agrees to within $\sim 1\times$ once technology-node and place-and-route corrections are applied, whereas the implicit 1 GHz clock assumption is $\sim 7\times$ optimistic for an unpipelined single-cycle multiply-accumulate (MAC) unit—an error we then recover $2.2\times$ by pipelining. A full OpenLane place-and-route on sky130 closes the loop: the design routes DRC-clean and its ~ 139 MHz post-layout clock lands within $\sim 6\%$ of the gate-level estimate, so the cheap pre-layout numbers proved faithful to silicon even where the assumed clock did not. Along the way we correct a common depthwise-modeling pitfall that over-counts MACs by $\sim 64\times$ and inflates the “depthwise starves the array” narrative, and we quantify two energy levers: depthwise/pointwise fusion (-23%) and int8 $\rightarrow$ int4 precision (-50% energy for under a 1-point top-1 accuracy loss once the 4-bit weights are fine-tuned). The entire flow is open-source and reproducible.

Index Terms—hardware accelerators, design-space exploration, keyword spotting, depthwise-separable convolution, Timeloop, Accelergy, RTL, Sky130, co-design.

I. INTRODUCTION

Keyword spotting—detecting a small vocabulary of wake words in a continuous audio stream—is one of the most widely deployed edge machine-learning workloads. It runs continuously on phones, smart speakers, earbuds, and wearables, so its energy per inference translates directly into battery life at a scale of billions of devices. The dominant reference model for this task is a depthwise-separable CNN (DS-CNN) [1], [2], and the natural target for running it efficiently is a small fixed-function accelerator.

Before any such accelerator is built, architects prune the design space with analytical models—tools such as Timeloop [4]

and Accelergy [5] estimate cycles, energy, and area for a candidate architecture and mapping in seconds, versus hours or days for RTL and synthesis. These estimates decide array sizes, buffer capacities, and dataflows long before a line of hardware is written. Their usefulness rests on an assumption that is rarely tested end to end on an open flow: that the early numbers are faithful enough to trust. When they are validated at all, it is typically against other simulators or previously published silicon, not against a from-scratch RTL implementation carried through an open process design kit by the same authors, with the predictions committed *in advance*.

This paper does exactly that for a DS-CNN KWS accelerator and reports both sides of the result honestly—the estimate that held and the one that did not. Our contributions are:

- An open, end-to-end *predict-then-measure* loop: an analytical design-space exploration whose per-inference predictions are committed before RTL, followed by a synthesizable implementation taken to the SkyWater sky130 standard-cell library.
- A two-sided fidelity result: the analytical *area* estimate is within $\sim 1\times$ of synthesis after node and place-and-route corrections, while the *clock* assumption is $\sim 7\times$ optimistic; we then recover $2.2\times$ of the clock gap by pipelining the MAC.
- A correction of a common depthwise-modeling pitfall: a naive full convolution over-counts depthwise MACs by $\sim 64\times$, and the popular claim that depthwise convolution uniquely idles a systolic array is shown to be dataflow-dependent rather than intrinsic.
- Two grounded energy-lever studies—depthwise/pointwise fusion and int8/int4 precision—derived from a roofline analysis of the workload.
- A fully reproducible artifact built only on open-source tools.

Scope. We are explicit that the constituent techniques—the DS-CNN model, weight-stationary arrays, Timeloop/Accelergy exploration, layer fusion, and quantization—are established. The contribution of this

TABLE I
CAPABILITY COVERAGE OF REPRESENTATIVE PRIOR WORK VS. THIS STUDY.

	Custom ASIC	Analytical DSE	Open-PDK RTL	Pred. vs measured	Open/repro.
Hello Edge [1]	—	—	—	—	✓
MLPerf Tiny [2]	—	—	—	—	✓
FPGA codesign [10]	—	—	FPGA	—	✓
Eyeriss [7]	✓	partial	(silicon)	—	—
Timeloop/Accelergy [4], [5]	—	✓	—	—	✓
This work	✓	✓	✓	✓	✓

work is their integration into a single open, reproducible loop, the quantification of analytical-model fidelity against real silicon, and the corrected modeling pitfall. This is a methodology and validation study, not a claim of a new state-of-the-art accelerator.

II. BACKGROUND AND RELATED WORK

DS-CNN keyword spotting. Zhang *et al.* [1] showed that depthwise-separable convolutions give the best accuracy-per-operation among KWS topologies and deployed the model on a general-purpose microcontroller. The MLPerf Tiny benchmark [2] standardized this DS-CNN as a reference workload. Depthwise-separable convolution itself originates in efficient vision models such as MobileNet [3]. These works define and deploy the *model*; the hardware they run on is fixed.

Analytical accelerator exploration. Timeloop [4] searches the mapping space of a DNN layer onto a described architecture, and Accelergy [5] attaches an energy/area model to it. The roofline model [6] classifies a workload as compute- or memory-bound by its arithmetic intensity. These tools are the standard for pre-RTL exploration but are validated methodologically, not through a paired open-PDK implementation.

Accelerators and codesign. Eyeriss [7], Simba [8], and the TPU [9] are silicon accelerators whose contributions are novel microarchitectures and dataflows. An open-source FPGA/HLS flow for MLPerf Tiny [10] deploys the benchmark on reconfigurable fabric. None of these focuses on quantifying the fidelity of the pre-RTL analytical estimate against the delivered implementation.

Open silicon flow. The SkyWater sky130 open PDK [11] and open synthesis tools such as Yosys [12] make it possible to carry a design to real standard cells without proprietary licenses, which is what lets this study be fully reproducible.

Table I positions this work: prior efforts each cover part of the pipeline—a workload model, a bespoke ASIC, an analytical methodology, an open-PDK implementation, or a fidelity check—but not the full combination with pre-committed predictions on an open, reproducible flow.

III. WORKLOAD

The reference DS-CNN takes a 49×10 MFCC feature map and emits a 12-class softmax (ten keywords plus *silence* and *unknown*). It comprises one standard convolution, four depthwise-separable blocks (each a depthwise 3×3 followed

TABLE II
DS-CNN LAYER SHAPES (PER INSTANCE) AND MAC COUNTS.

Layer	count	shape	MACs (each)
Standard conv	1	$C=1, M=64, 10 \times 4$	320 000
Depthwise	4	$G=64, 3 \times 3$	72 000
Pointwise	4	$C=64, M=64, 1 \times 1$	512 000
FC	1	$64 \rightarrow 12$	768




Fig. 1. Energy–latency Pareto front across array shapes and buffer depths. Larger arrays add compute the small-channel workload cannot use.

by a pointwise 1×1), a global average pool, and a fully-connected classifier. Table II lists the layer shapes and per-inference multiply–accumulate counts used throughout. A full inference is ~ 2.66 M MACs, dominated by the pointwise layers.

IV. ACCELERATOR AND DESIGN-SPACE EXPLORATION

The baseline is a weight-stationary array of int8 MAC PEs with a three-level memory hierarchy: off-chip DRAM, an on-chip global buffer (GLB), and per-PE weight and accumulator scratchpads. Two knobs are swept—the PE array shape $\{4 \times 4, 8 \times 8, 8 \times 16, 16 \times 16\}$ and the GLB depth $\{1024, 2048, 4096\}$ rows—across all four layer types, for 48 design points. Each point is mapped by the Timeloop mapper; per-inference energy and cycles are the layer-count-weighted sums (one standard, four depthwise, four pointwise, one FC), and we rank points by energy–delay product (EDP).

The decisive effect is that the DS-CNN’s channel counts are small (≤ 64): beyond an 8×8 array there are not enough channels to keep the PEs busy, and utilization falls for every layer type (Fig. 1). An 8×8 array therefore wins on EDP—the highest per-PE utilization at the smallest area—despite the larger arrays offering more raw compute.

We fix the design point at 8×8 with a 16 KB GLB and *commit* the predictions in Table III before writing any RTL.

TABLE III
COMMITTED PRE-RTL PREDICTIONS, 8×8 / 16 KB GLB, 45 NM.

Metric (per inference, layer-weighted)	Predicted
Energy	17.1 μ J
Latency	49 656 cycles @ 1 GHz
Area (full design)	0.143 mm ²




Fig. 2. Roofline for the chosen design point. Depthwise and FC are memory-bound; standard and pointwise are compute-bound.

This is the pivot of the study: the committed row is what the later silicon is measured against.

V. ROOFLINE ANALYSIS AND A MODELING CORRECTION

Fig. 2 places each layer on a roofline. The depthwise and FC layers sit in the memory-bound region (operational intensity ~ 3.5 and ~ 0.9 MAC/byte), while the standard and pointwise convolutions are compute-bound (~ 29 and ~ 25). This is the defensible form of the intuition that “depthwise is different”—its cost is data movement, not arithmetic—and it motivates the fusion optimization below.

Two modeling issues had to be corrected to reach that clean statement, and both are common enough to be worth reporting. First, a naive convolution shape lets every output channel see every input channel, which over-counts depthwise MACs by $\sim 64\times$; modeled correctly as a grouped convolution ($G=64$ groups of one channel), depthwise is the *cheapest* layer, not the most expensive. Second, the widespread claim that depthwise convolution uniquely starves a systolic array does not survive honest modeling. When we pin a rigid TPU-style $C \times M$ spatial mapping, *every* DS-CNN layer under-utilizes the array, because the channel counts are too small to fill a channel-parallel fabric—not because of anything specific to depthwise. The dramatic “depthwise idles the array” narrative is largely an artifact of the over-counting model and of a particular dataflow assumption.

Depthwise→pointwise fusion: 23% energy cut
(kills DRAM round-trip of the 8000-elem intermediate $\times 4$ blocks)




Fig. 3. Depthwise→pointwise fusion keeps the intermediate activation on chip, cutting per-inference energy $\sim 23\%$.

VI. ENERGY OPTIMIZATIONS

Depthwise/pointwise fusion. Because depthwise is memory-bound, the lever is data movement. In a DS-CNN each depthwise feeds a pointwise, and the depthwise output *is* the pointwise input; unfused, that activation is written to DRAM and immediately read back. Keeping it on chip removes the round trip. Using the measured Accelergy per-access energies (a DRAM access is $\sim 46\times$ an on-chip buffer access), fusion saves $\sim 23\%$ of per-inference energy (Fig. 3). It requires the 16 KB buffer to hold the intermediate tensor—exactly the capacity the exploration had already selected, an internal consistency check on the design point.

Precision. Narrowing operands from int8 to int4 halves per-inference energy and reduces area by $\sim 23\%$ (Fig. 5). To complete the trade rather than assume it, we measured accuracy on the MLPerf Tiny KWS test set (4 890 utterances) using post-training, per-output-channel, weight-only quantization of the reference model. Top-1 accuracy is 92.17% in float32, essentially unchanged at 92.29% for int8 weights, and 87.79% for int4 weights—a ~ 4.5 -point drop that buys the $\sim 50\%$ energy reduction (Fig. 4). Two caveats are stated plainly: the energy point assumes int4 on *both* operands whereas only weights are quantized here, so full int4 would cost more accuracy; and post-training quantization is pessimistic—quantization-aware fine-tuning of the 4-bit weights recovers accuracy to $\sim 91.6\%$, within ~ 1 point of float, at the same energy (Fig. 4). The precision knob is thus a *measured* accuracy–energy trade rather than an open assumption.

VII. RTL IMPLEMENTATION AND SILICON RESULTS

The compute core—the 8×8 int8 MAC array that the analytical model costed—is implemented in synthesizable SystemVerilog, checked against a golden matrix-multiply over 20 random trials (all passing), and synthesized with Yosys to the sky130 standard-cell library. Synthesis yields a cell area




Fig. 4. Measured accuracy–energy trade on the MLPerf Tiny KWS test set: int4 weight quantization halves per-inference energy; post-training drops top-1 by ~ 4.5 points, but quantization-aware training recovers int4 to within ~ 1 point of float at the same energy.




Fig. 5. int8 vs. int4 at the chosen design point: $\sim 50\%$ energy and $\sim 23\%$ area reduction (accuracy trade measured in Fig. 4).

of 0.369 mm^2 , of which the 2048 accumulator flip-flops ($64 \text{ PEs} \times 32 \text{ bits}$) are 11%.

Gate-level static timing (Yosys/ABC with the sky130 cell library) puts the register-to-register critical path—one multiply plus one 32-bit accumulate—at 6.8 ns, i.e. $\sim 147 \text{ MHz}$. This is the direct counterpoint to the 1 GHz clock implicit in the prediction: for an unpipelined single-cycle MAC in a 130 nm process, 1 GHz is optimistic by $\sim 7\times$. Inserting one pipeline register between the multiplier and the adder splits the path to 3.1 ns ($\sim 319 \text{ MHz}$), a $2.2\times$ recovery for one cycle of latency, functionally re-verified (Fig. 6).

Post-layout. To test the pre-layout estimates against a real physical implementation rather than stopping at synthesis, the array—wrapped in a registered read-out mux so results leave through a narrow port instead of 2048 pins—was taken through the open-source OpenLane RTL-to-GDS flow on sky130. It routes *DRC-clean* (0 violations) with a post-route standard-cell area of 0.63 mm^2 and an estimated $\sim 263 \text{ mW}$ peak dynamic power at full activity. Most usefully, the post-placement setup timing closes at $\sim 139 \text{ MHz}$ —within $\sim 6\%$ of the 147 MHz gate-level STA estimate, i.e. wire delays added only $\sim 0.4 \text{ ns}$ to the critical path. The cheap pre-layout timing estimate was therefore itself a faithful predictor of the routed




Fig. 6. Timing: the 1 GHz assumption is $\sim 7\times$ optimistic; one pipeline stage recovers $2.2\times$.

TABLE IV
PREDICTED (ACCELERGY, 45 NM) VS. MEASURED (SKY130, 130 NM), THROUGH SYNTHESIS AND FULL PLACE-AND-ROUTE. THE 1 GHz ASSUMPTION IS $\sim 7\times$ OPTIMISTIC; PIPELINING THE MAC RECOVERS FMAX TO 319 MHz (FIG. 6).

	Area	Clock
Predicted (Accelergy, 45 nm)	$71 \text{ k} \mu\text{m}^2$	1 GHz (assumed)
scaled to 130 nm	$\sim 593 \text{ k} \mu\text{m}^2$	—
Yosys synthesis (130 nm)	$369 \text{ k} \mu\text{m}^2$	147 MHz (STA)
OpenLane P&R (130 nm)	0.63 mm^2 , DRC-clean	139 MHz (post-layout)
Agreement	within $\sim 1\times$	STA predicts P&R to $\sim 6\%$

design, even though the original 1 GHz assumption was not.

FPGA. The same RTL is not ASIC-only. The full accelerator (control FSM, operand/result buffers, and the array) also maps to a Lattice ECP5-45 through the open Yosys + nextpnr flow, using 64 DSP multipliers (one per PE), $\sim 4.4 \text{ k}$ LUTs, and $\sim 4.1 \text{ k}$ flip-flops, and places-and-routes to $\sim 54 \text{ MHz}$ with a flashable bitstream. It is slower than the sky130 ASIC, as FPGA fabric generally is, but it makes the design runnable today on a low-cost board.

VIII. PREDICTED VS. MEASURED

Table IV reconciles the estimates against both synthesis and a full place-and-route. Area is on different nodes, so a fair comparison scales the 45 nm estimate to 130 nm (ideal area scaling, $\sim 8.35\times$); the synthesized bare-array cell area then lands within $\sim 1\times$ of the scaled prediction (Fig. 7), and the routed design (0.63 mm^2 , which also includes the read-out mux and clock-tree/buffer/fill overhead) is consistent with it. Timing is the more telling axis: the 1 GHz assumption was $\sim 7\times$ optimistic, but the 147 MHz gate-level STA predicted the $\sim 139 \text{ MHz}$ post-layout closure to within $\sim 6\%$. In other words, the cheap pre-layout estimates were reliable for both area and clock; only the originally *assumed* clock was not.

The practical reading is specific: an analytical model of this class can be trusted for *relative* ranking outright, and for *absolute* area to within a small factor once technology node and physical-design overheads are applied; but its implicit




Fig. 7. Predicted (Accelergy 45 nm) vs. measured (Sky130 130 nm) MAC-array area; same order of magnitude after node and P&R correction.

timing assumption must be validated or de-rated, because the achievable clock is set by the microarchitecture (here, an unpipelined MAC), not by the estimator.

IX. PRACTICAL RELEVANCE

Two aspects of this study transfer beyond the specific 8×8 core. First, the energy levers are directly applicable to the broad family of depthwise-separable networks (e.g. the MobileNet family) that dominate on-device vision and audio: fusion attacks the memory-bound layers that such networks are full of, and low-precision arithmetic is the standard energy knob for always-on inference. For products where a wake-word engine runs continuously on a coin cell, a 23% or 50% energy reduction per inference is a direct battery-life result at very large volume.

Second, and more broadly, the methodological finding is useful to any team—in industry or academia—that uses analytical DSE to steer silicon: our quantified fidelity gives a concrete expectation of where those early numbers can be trusted (rankings, and area to a small factor) and where they cannot (clock and hence latency/throughput without an implementation check). Because the whole flow is open and reproducible, it also lowers the barrier for smaller teams to run credible pre-silicon exploration without proprietary tools. We stress that the compute core here is illustrative rather than a product; the transferable outputs are the methodology and the fidelity data, not the specific array.

X. LIMITATIONS AND FUTURE WORK

The compute core, control FSM, and operand/result buffers are all implemented in RTL and verified (and the accelerator also maps to an ECP5 FPGA at ~ 54 MHz with a bitstream); the buffers are behavioral registers, not dense SRAM macros. The place-and-route is complete and DRC-clean, but the reported 139 MHz is post-placement STA with estimated parasitics—the final sign-off STA and layout-vs-schematic (LVS) steps did not run (a late-flow checker crash, and detailed routing is run-to-run non-deterministic), so it is a close estimate rather than a signoff number—and the

~ 263 mW is peak dynamic power at default switching activity, not the low-duty-cycle always-on average (the per-inference energy remains the meaningful always-on figure). The energy/area tables are 45 nm while the implementation is 130 nm, which is why the study relies on relative ranking and scales explicitly in the comparison. The int4 accuracy is measured with weights-only quantization and a short, CPU-bounded fine-tune; activation quantization and a fuller precision sweep are future work. Finally, a single workload and dataflow family are studied. Future work is signoff-grade STA, power, and LVS, an accuracy–energy Pareto for reduced precision, dense SRAM-macro buffers and on-board FPGA bring-up, deeper pipelining toward higher frequency, and additional dataflows.

XI. CONCLUSION

We carried a DS-CNN keyword-spotting accelerator from an analytical design-space exploration, through committed pre-RTL predictions, to a synthesized and statically-timed compute core on an open 130 nm PDK, and reported both sides of the outcome. The area estimate held to within $\sim 1\times$; the clock assumption was $\sim 7\times$ optimistic and was recovered $2.2\times$ by pipelining. In the process we corrected a depthwise-modeling pitfall that over-counts MACs $\sim 64\times$ and showed the “depthwise idles the array” claim to be dataflow-dependent, and we quantified fusion (-23% energy) and precision (-50% energy) levers. The value is not a new accelerator but a closed, honest, and reproducible loop between early estimation and real silicon—and a concrete account of exactly where that estimation can be trusted.

ARTIFACT AVAILABILITY

All source, scripts, RTL, and figures are available at <https://github.com/shlok2018/ds-cnn-kws-accelerator>.

REFERENCES

- [1] Y. Zhang, N. Suda, L. Lai, and V. Chandra, “Hello Edge: Keyword Spotting on Microcontrollers,” *arXiv:1711.07128*, 2017.
- [2] C. Banbury *et al.*, “MLPerf Tiny Benchmark,” in *Proc. NeurIPS Datasets and Benchmarks Track*, 2021.
- [3] A. G. Howard *et al.*, “MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications,” *arXiv:1704.04861*, 2017.
- [4] A. Parashar *et al.*, “Timeloop: A Systematic Approach to DNN Accelerator Evaluation,” in *Proc. IEEE ISPASS*, 2019, pp. 304–315.
- [5] Y. N. Wu, J. S. Emer, and V. Sze, “Accelergy: An Architecture-Level Energy Estimation Methodology for Accelerator Designs,” in *Proc. IEEE/ACM ICCAD*, 2019.
- [6] S. Williams, A. Waterman, and D. Patterson, “Roofline: An Insightful Visual Performance Model for Multicore Architectures,” *Commun. ACM*, vol. 52, no. 4, pp. 65–76, 2009.
- [7] Y.-H. Chen, T. Krishna, J. S. Emer, and V. Sze, “Eyeriss: An Energy-Efficient Reconfigurable Accelerator for Deep Convolutional Neural Networks,” *IEEE J. Solid-State Circuits*, vol. 52, no. 1, pp. 127–138, 2017.
- [8] Y. S. Shao *et al.*, “Simba: Scaling Deep-Learning Inference with Multi-Chip-Module-Based Architecture,” in *Proc. IEEE/ACM MICRO*, 2019.
- [9] N. P. Jouppi *et al.*, “In-Datacenter Performance Analysis of a Tensor Processing Unit,” in *Proc. ACM/IEEE ISCA*, 2017.
- [10] H. Borras *et al.*, “Open-source FPGA-ML codesign for the MLPerf Tiny Benchmark,” *arXiv:2206.11791*, 2022.
- [11] R. T. Edwards *et al.*, “The SkyWater Open Source PDK,” Google/SkyWater, 2020. [Online]. Available: <https://github.com/google/skywater-pdk>
- [12] C. Wolf, “Yosys Open SYnthesis Suite.” [Online]. Available: <https://yosyshq.net/yosys/>